

VMix record ingest — near-real-time, Nextcloud-aware

Goal: same as `docs/atem-iso-ingest.md` — get each VMix PC's growing recording onto the mothership's Nextcloud as close to real time as possible, without corrupting anything and without falling permanently behind. Covers all 4 VMix PCs: VMix Node 1 (`192.168.20.21 / .22` , off Theatre 1) and VMix Node 2 (`192.168.21.21 / .22` , off Theatre 4) — see `docs/ip-address-map.md` .

Why this is actually simpler than the ATEM case

The ATEM Mini Extreme ISO only exposes its recording drive over FTP — a file-transfer protocol, not a filesystem-access protocol, so a mount-based approach there needs `rclone mount` as a FUSE bridge (that's the ATEM doc's *alternative* path; its default pulls via FTP `REST` resume with no mount at all). **VMix PCs run Windows, which natively serves SMB** — a real network filesystem protocol, designed to be mounted. There's no bridging trick needed here; SMB is the native way two machines share a filesystem on a LAN. That makes this the more natural fit for a mount-based approach, not a stretch of the ATEM pattern onto a different protocol.

Architecture

Same shape as the ATEM ingest, reusing the identical tools (`config/vmix-record-ingest/` mirrors `config/atem-iso-ingest/rclone-mount-rsync/`), as its own Unraid container (`192.168.1.18`):

```
VMix PC (192.168.2X.2X, SMB share) --[Tailscale, via that VMix node's own Slate AX]--> Unraid
```



`rsync --append` (not rsync's general checksum-diff algorithm) is what makes this incremental — it seeks straight to the destination's current size and transfers only the new tail, rather than reading/hashing the whole file. Same reasoning as the ATEM ingest's `rclone-mount-rsync` alternative; see that doc for the full explanation of why this differs from rsync's usual delta-transfer mode. (Checked with rsync 3.2.7: on

a growing file, a repeat `--append` pass sends only the new tail and the sender seeks straight to the old size. Two limits of the same mode: a file whose earlier bytes are rewritten in place — e.g. an MP4 header patched when recording stops — is never corrected by `--append`, and `--append-verify` *skips* any file that is already the same size on both sides, so it is not a whole-file integrity check. `VERIFY=1` therefore runs a full `--checksum --inplace` pass instead — see the script.)

SMB, like FTP, has no change notification in rclone, so a growing file's size on the mount only refreshes when rclone's directory cache expires — default 5 minutes, which would quietly make the 90s sync interval a ~5-minute one. The script sets `--dir-cache-time 30s` (rclone mount docs).

Second write destination for live editing. Same as the ATEM ingest: the pull also writes to the edit suite's dedicated NAS (`192.168.22.x`) alongside Nextcloud's External Storage, one read off the VMix PC's SMB share, two writes. See `docs/live-editing.md` for the editing subsystem this feeds — not yet implemented in `config/vmix-record-ingest/mount-and-sync.sh`, tracked in `docs/open-questions.md`.

A native Linux CIFS mount (`mount -t cifs`) is a simpler alternative to `rclone mount` for SMB specifically — no FUSE bridge needed at all, since the kernel has first-class SMB client support. The tooling here uses rclone anyway, to reuse the exact same config/script pattern already built for the ATEM ingest rather than introduce a third mechanism. Worth reconsidering if operational simplicity outweighs consistency.

Why bandwidth isn't a shared concern with the ATEM ingest

Each VMix node has **its own dedicated GL-iNet Slate AX (GL-AXT1800) uplink and its own Tailscale connection** — it is not part of the theatre subnet it physically sits near (see `docs/topology.md`). So VMix Node 1's record-ingest traffic rides over Node 1's own link, entirely separate from Theatre 1's ATEM ISO ingest and SRT traffic. There's no contention between the two ingest pipelines sharing a link — each just needs to fit inside its own node's ceiling — an estimated ~150–250 Mbps combined under Tailscale, not yet benchmarked (`docs/open-questions.md` #19).

Open items

- **Confirm what VMix is actually recording** — program mix only, or per-input ISO the same way the ATEMs do — and at what resolution/codecs/bitrate. None of this is assumed here; it directly determines the real bandwidth load on each VMix node's uplink. (Note `docs/server-specification.md` does pencil in ~30 Mbps per PC for pool sizing — an explicit placeholder, not a confirmed figure.)
- **Set up SMB sharing on the real PCs** — confirm the actual recording folder path and share name, and use a dedicated, minimal-privilege (read-only) local account for the ingest process rather than an admin/shared login.
- **Same partial-file-safety question as the ATEM ingest, but easier to answer:** the answer depends entirely on which container VMix is set to record (a plain MP4 with its index written only at the end is not readable until recording stops; a fragmented MP4 or a streaming-friendly container is). Not established here — check empirically (open a currently-recording file's ingested copy in a player), and check the *finished* copy too, since an append-only mirror can miss a header rewritten at stop (see above).

- **Confirm the version-retention setting on Nextcloud** covers these folders too (see the same caveat in `docs/atem-iso-ingest.md`) — same risk of accumulating snapshots on a repeatedly-changing external-storage file.